



THE UNIVERSITY
OF QUEENSLAND
A U S T R A L I A

Auditing VLM Perception of
Gender Bias using Human Sexism
Scales in Image Search

by

GIA HUNG HUYNH

*under the supervision of
Prof Gianluca Demartini*

Use of AI Statement

I acknowledge the use of generative AI tools in completing this assessment. Details of which tools were used and how they were used are provided in the table below, along with appropriate in-text and full references. I take responsibility for critically evaluating and integrating the AI-generated content, and ensuring it adheres to academic integrity standards.

Use of AI in this project proposal

- **Scope:** This declaration applies to this project proposal submission only.
- **Research question refinement:** ChatGPT-4o and Gemini were used to test alternative phrasing and improve precision of the research question.
- **Literature review support:** NotebookLM, Gemini, and ChatGPT-4o were used to summarize and compare relevant papers; final interpretation and integration were independently verified and rewritten by the author.
- **Language editing:** ChatGPT-4o, Gemini, and Cursor AI were used to improve grammar, style, and clarity of academic writing.
- **Content creation and formatting:** Cursor AI was used to assist with LaTeX code editing, structural section updates, and formatting tasks (including table/layout adjustments). Gemini and ChatGPT-4o were also used to draft/revise text snippets that were critically reviewed before inclusion.
- **Accountability:** All final decisions, technical judgments, and submitted content remain the author's responsibility.

AI Model and date	Table 1. AI tool usage (tick all that apply)								
	Language Translation	Grammar / style / spelling	Topic exploration	Research question	Literature review	Content Creation visual (formatting)	Content Creation (e.g. code, text)	Feedback	Other (provide details)
ChatGPT-4o 2026		✓		✓	✓		✓		
Cursor AI assistant 2026		✓				✓	✓		
Gemini 2026		✓		✓	✓		✓		
NotebookLM 2026					✓				Document synthesis

Abstract

Vision-Language Models (VLMs) are increasingly deployed as automated judges in content moderation and image retrieval pipelines, yet they systematically inherit gender stereotypes from their visual training distributions. These biases originate at the visual-semantic alignment stage, where vision encoders project images into a shared embedding space, and have been shown to skew zero-shot retrieval outcomes in ways that disadvantage women in professional contexts.

A critical barrier to detecting these biases is *Algorithmic Defensiveness*: when safety-aligned models—trained with RLHF and Constitutional AI—detect evaluative framing in an input, they enter a defensive mode and produce sanitized outputs that mask their latent distributional skew. Existing auditing methodologies rely on model self-report and are therefore unable to distinguish genuine debiasing from defensive evasion, leaving a fundamental gap in the AI auditing literature.

This project addresses that gap by introducing a human-anchored auditing framework built around implicit logit extraction. Rather than analysing generated text, the pipeline captures raw token probability distributions at the output layer prior to autoregressive sampling, bypassing the safety overlay entirely. These logit distributions are mapped against the Ambivalent Sexism Inventory (ASI) to produce a quantified distributional bias profile for each model. The framework operates in two phases: Phase 1 establishes an intrinsic baseline by administering ASI items through five implicit prompt structures across three modality conditions; Phase 2 evaluates extrinsic downstream behavior by deploying the same models on historical and contemporary image search queries drawn from the SIGIR 2018 dataset, benchmarked against high-ASI and low-ASI human reference groups. The CALM positional option-shuffling framework is applied throughout to isolate the vision encoder’s specific contribution to biased judgments from linguistic confounds, quantified via the robustness rate differential ΔRR . Three state-of-the-art VLMs—Llama-3.2-Vision (90B), GPT-4o, and Qwen2-VL (72B)—are evaluated across 3,150 model inferences. By grounding VLM audit outcomes in established human psychometric instruments and real-world retrieval data, this research provides the scientific community with a trustworthy, replicable methodology for exposing hidden representational bias in multimodal AI systems.

Contents

Use of AI Statement	ii
Abstract	iv
List of Figures	vii
List of Tables	viii
1 Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	1
1.3 Scope and Significance	2
2 Background and Literature Review	3
2.1 Background and Theory	3
2.1.1 Architecture: The Dual-Encoder Vulnerability	3
2.1.2 Quantifying Bias: The Ambivalent Sexism Inventory (ASI)	3
2.1.3 Auditing Robustness: The CALM Framework and Defensive Evasion	3
2.2 Literature Review	4
2.2.1 The "Biased Eyes" Phenomenon	4
2.2.2 The Vulnerabilities of Automated Judges	4
2.2.3 Algorithmic Defensiveness and Testing Mode	4
2.3 Literature Gap	5
2.3.1 The Auditing Deadlock	5
2.3.2 The Absence of Human-Centric Grounding	5
2.3.3 Proposed Novelty: Implicit Logit Mapping	5
3 Project Plan	7
3.1 Project Aims	7
3.2 Methodology	8
3.2.1 Agile Pipeline Engineering and System Architecture	8
3.2.2 Phase 1: Intrinsic Baseline Measurement & Modality Activation	8
3.2.3 Phase 2: Extrinsic Downstream Evaluation	10
3.2.4 Modality Isolation via the CALM Framework	10
3.2.5 Statistical Analysis and Evaluation Metrics	11

3.2.6	Experimental Scope and Analysis Parameters	12
3.3	Milestones and Timelines	13
3.4	Health and Safety and Risk Assessments	13
3.5	Ethics statement	14

List of Figures

List of Tables

3.1	Summary of Analysis Parameters.	12
3.2	Planned research phases (April–November 2026).	13

1. Introduction

1.1 Background and Motivation

The rapid integration of Vision-Language Models (VLMs) into critical moderation and retrieval systems demands a shift toward robust, objective auditing protocols. Without these mechanisms, algorithmic biases risk proliferating unchecked across digital platforms. The adoption of the "LLM-as-a-Judge" paradigm for content moderation is increasing rapidly; however, modern Vision-Language Models (VLMs)—including Llama-3.2-Vision [1], GPT-4o, and Qwen2-VL [2]—frequently inherit deep-seated gender stereotypes from their visual training distributions [3, 4]. Because these generative models rely on vision encoder modules to project images into a shared semantic space with text embeddings [5], they exhibit pronounced gender and occupational biases that originate at the visual-semantic alignment stage. The specific encoder architecture varies across models (e.g., CLIP-derived for Llama-3.2-Vision; proprietary ViT-based for Qwen2-VL), and architectural attribution is treated as exploratory within this study. Research has demonstrated that this "Biased Eyes" phenomenon actively skews zero-shot retrieval outcomes [6]. Consequently, these automated systems may perpetuate gender bias in downstream retrieval contexts [7].

1.2 Problem Statement

Despite these identified risks, standard AI safety audits often fail to detect such prejudices. This failure is rooted in *Algorithmic Defensiveness*—a systematic pattern, identified in this study, whereby safety-aligned models produce sanitized outputs when evaluative framing is detected in the input. Modern models are optimized with stringent safety alignments, such as Reinforcement Learning from Human Feedback (RLHF) and Constitutional AI [8, 9]. When presented with overt evaluative framing, models may enter a "testing mode," generating artificially sanitized and defensive responses to satisfy internal safety constraints [10, 11]. Furthermore, the LLM-as-a-Judge framework is susceptible to intrinsic cognitive biases, including positional and self-enhancement effects [12, 13, 14]. Because AI-based evaluators exhibit both inherent bias and defensive evasion, relying on them as the primary auditing mechanism introduces significant methodological vulnerabilities. There is a distinct gap in the literature for an auditing methodology that can successfully bypass these defensive filters to measure the latent distributional biases in multimodal models.

1.3 Scope and Significance

To address these flaws, this project introduces a human-anchored auditing framework that maps VLM zero-shot judgments against Ambivalent Sexism Inventory (ASI) baselines [15] grounded in real-user data. This project asks whether, when a safety-aligned VLM performs implicit evaluation of image search results, bias in the vision encoder dominates higher-level aligned reasoning. It further tests whether resulting judgments align with human high-ASI cognitive patterns. To systematically answer this, the research is structured to address two distinct, sequential questions:

1. What is the intrinsic psychological baseline attitude of the AI when evaluated in a vacuum?
2. Is that intrinsic distributional skew descriptively consistent with the model’s extrinsic, downstream behavior when evaluating real-world search queries?

The scope of this research is based on the historical SIGIR 2018 dataset [16] to evaluate real human perception, using high-ASI and low-ASI reference-group baselines. To bypass safety alignments, the project will employ implicit tasking, which utilizes indirect, objective grading scales (logit distributions) rather than direct questioning. Furthermore, the study will integrate the CALM evaluation framework [12] to isolate the origin of the biases. By demonstrating that implicit tasking and human psychological scales can effectively bypass defensive filters, this research advances the science of AI auditing and provides the scientific community with a trustworthy methodology for detecting hidden biases in multimodal models.

2. Background and Literature Review

2.1 Background and Theory

2.1.1 Architecture: The Dual-Encoder Vulnerability

Contemporary VLMs such as Llama-3.2-Vision [1] and Qwen2-VL [2] rely on dedicated vision encoder modules to project images into a shared semantic space alongside text embeddings [5]. While the specific encoder architecture differs across models—CLIP-derived contrastive encoders in some cases, proprietary ViT-based modules in others—all share the structural characteristic that visual-semantic associations are formed prior to higher-level language reasoning or safety-alignment layers. Consequently, the visual pathway can preserve distributional bias signals that aligned language-layer behaviour may not fully suppress.

2.1.2 Quantifying Bias: The Ambivalent Sexism Inventory (ASI)

To move beyond subjective observation, this research maps VLM token-probability distributions against the Ambivalent Sexism Inventory (ASI) [15]. As a foundational psychological framework, the ASI splits sexist attitudes into two categories: Hostile Sexism (antagonistic) and Benevolent Sexism (patronizing and restrictive). By using this established human scale, we obtain a validated reference instrument to measure whether a VLM’s output distributions exhibit patterns consistent with human sexist attitudes.

2.1.3 Auditing Robustness: The CALM Framework and Defensive Evasion

While the industry has pivoted toward the "LLM-as-a-Judge" framework for automated moderation, these systems are not neutral observers. To expose the structural flaws and "defensive evasion" common in safety-aligned models, this study integrates the CALM evaluation framework [12]. CALM functions as an automated "attack-and-detect" methodology, systematically perturbing inputs to evaluate judgment stability. By applying these systematic perturbations, we can bypass the sanitized

”testing mode” and reveal the underlying distributional biases within the evaluator.

2.2 Literature Review

2.2.1 The ”Biased Eyes” Phenomenon

Foundational Vision-Language Models carry inherent visual vulnerabilities. Jiang et al. [3] and Weng et al. [4] demonstrated that within multimodal architectures, image features contribute significantly more to stereotypical bias than text features, effectively allowing the visual modality to override linguistic reasoning. Ghate et al. [6] established that these intrinsic representational biases do not remain isolated; they systematically propagate into zero-shot retrieval outcomes. This propagation indicates that the visual encoder can systematically distort model judgments related to gender and occupation.

2.2.2 The Vulnerabilities of Automated Judges

To scale moderation, the industry heavily relies on the ”LLM-as-a-Judge” paradigm. Yet, utilizing a biased system to audit another biased system introduces severe methodological flaws. Chen et al. [13] and Li et al. [14] demonstrated that automated evaluators suffer from distinct cognitive vulnerabilities, while Ye et al. [12] highlighted their susceptibility to positional, egocentric, and bandwagon biases. Without rigorous perturbation testing, LLM-as-a-Judge frameworks risk creating a compounding cycle of unverified bias.

2.2.3 Algorithmic Defensiveness and Testing Mode

Standard text-based prompting struggles to detect these deep-seated biases because modern AI systems exhibit active evasion. Gao and Kreiss [10] and Yang et al. [11] identified the critical flaw of overt evaluative framing: when a prompt hints that a model is being tested on social reasoning, it triggers a defensive ”testing mode.” Triggered by evaluative framing, safety-alignment mechanisms incentivize the generation of sanitized outputs that obscure underlying visual-semantic associations. As a result, direct evaluative prompting is insufficient for detecting latent visual prejudices.

2.3 Literature Gap

2.3.1 The Auditing Deadlock

The existing literature presents a significant auditing deadlock. While it is established that visual encoders actively propagate gender bias into downstream tasks [6], standard evaluation methods fail to detect these biases because direct, overt prompting triggers defensive safety alignments [10]. Furthermore, the reliance on other AI models as judges introduces secondary, unverified biases, rendering the current "LLM-as-a-Judge" paradigm functionally insufficient for objective safety audits. Consequently, there is a critical gap for a methodology that can extract a model's true visual associations without alerting its defensive filters.

2.3.2 The Absence of Human-Centric Grounding

A primary limitation in current VLM research is the lack of human-anchored grounding. Most AI auditing frameworks remain self-referential, utilizing AI-generated benchmarks to evaluate AI behavior. While the historical SIGIR 2018 dataset [16] provides a rigorous baseline for how sexist and objective humans perceive bias in image search, this psychological data has not been utilized as a "gold standard" to measure the latent profiles of generative VLMs. There is no existing framework that maps an AI's zero-shot visual judgments directly against established human psychological scales like the Ambivalent Sexism Inventory (ASI) [15].

2.3.3 Proposed Novelty: Implicit Logit Mapping

To bridge these gaps, this research introduces an auditing framework that bypasses textual safety filters via implicit tasking and logit-level extraction. By capturing raw probability distributions prior to token generation, the methodology derives the model's intrinsic bias profile ($ASI_{intrinsic}$) directly from its internal representational state. Crucially, this $ASI_{intrinsic}$ is a model-specific metric calculated independently of any human participant scores. This baseline is then benchmarked against the model's downstream evaluations on the SIGIR 2018 query set, compared to historical high-ASI and low-ASI human reference groups. This comparison tests whether the model's internal stereotypes align with established human bias patterns. This separation between model-side intrinsic estimation and human-anchored extrinsic benchmarking provides an objective audit protocol that is currently absent from the generative VLM literature.

This project addresses this gap by mapping raw VLM logit distributions to a model-intrinsic $ASI_{intrinsic}$ signal, then benchmarking downstream behavior against human-validated high-ASI/low-ASI reference groups, establishing a more transparent and objective auditing baseline than current text-based AI judges.

3. Project Plan

3.1 Project Aims

The primary objective of this project is to investigate whether the intrinsic gender biases embedded within the visual encoders of state-of-the-art Vision-Language Models (VLMs) dictate their extrinsic behavior when evaluating real-world image search queries. To achieve this, the project will systematically map VLM outputs to the psychological Ambivalent Sexism Inventory (ASI) using a novel implicit logit extraction pipeline.

The specific aims of this research are:

- 1. To measure the intrinsic psychological baseline of state-of-the-art VLMs:** By administering the Ambivalent Sexism Inventory (ASI) items through five implicit prompt structures adapted from the syntactic perturbation methodology of Gao and Kreiss [10], and extracting raw logit distributions, this project aims to successfully bypass textual safety alignments and characterise the models’ intrinsic token-probability distributions.
- 2. To evaluate VLM performance on implicit downstream visual tasks:** Using the historical SIGIR 2018 image search dataset, this aim involves testing the models on real-world queries and mapping their automated judgements against established ”High-ASI” and ”Low-ASI” human baselines.
- 3. To characterise the consistency between intrinsic and extrinsic bias:** By computing rank-order comparisons between each model’s $ASI_{intrinsic}$ score and its extrinsic deviation from the human low-ASI baseline, this project characterises whether models with higher intrinsic bias scores produce correspondingly more biased downstream evaluations. Given the small model sample ($n = 3$), this comparison is treated as a descriptive case study; Spearman’s ρ is reported as a descriptive statistic only, and no cross-model inferential claim is made.
- 4. To validate that measured bias originates in the visual encoder rather than linguistic artifacts:** The CALM evaluation framework—applied through positional option-shuffling—serves as a methodological validation step to confirm that the bias signals detected in Aims 1 and 2 are attributable to the models’ visual encoders and not to prompt-level linguistic confounds.

3.2 Methodology

3.2.1 Agile Pipeline Engineering and System Architecture

This research will utilize an Agile development framework, initially engineering the evaluation pipeline around the open-weights Llama-3.2-Vision (90B) model [1] to establish a rapid testing baseline. Upon validation, the architecture will systematically scale to evaluate the closed-source GPT-4o (via API access) and the large-scale Qwen2-VL (72B) model [2]. The core system will be engineered in Python to directly interface with the models’ output layers. Rather than capturing generated text—which is highly susceptible to safety filter interference—the pipeline will extract raw logit probability distributions prior to token generation. This captures the models’ pre-sampling visual-semantic token probabilities and confidence scores. Additionally, to provide empirical proof of visual grounding, the pipeline will extract cross-attention weights to generate heatmaps, visually demonstrating which specific regions of an image trigger biased token generation.

3.2.2 Phase 1: Intrinsic Baseline Measurement & Modality Activation

Phase 1 establishes the intrinsic psychological baseline ($ASI_{intrinsic}$) for each model. To avoid the evaluative defenses triggered by direct grading, the Ambivalent Sexism Inventory (ASI) items are rephrased into five distinct linguistic structures adapted from the syntactic perturbation methodology of Gao and Kreiss [10]. This sensitivity analysis determines whether measured bias is a stable representational feature or a prompt-formulation artifact. The five structures include:

- **Direct Labeling:** *“Is this person too easily offended?”*
- **Linguistic Inversion:** *“Would it be incorrect to say this person is too easily offended?”*
- **Third-person Attribution:** *“Would an observer describe this person as too easily offended?”*
- **Hypothetical Scenario:** *“If this person were in a professional setting, would they be too easily offended?”*
- **Descriptive Template:** *“Can this person be described as too easily offended?”*

The methodology uses logit-level extraction to capture the probability distributions for binary response tokens (e.g., *yes* and *no*) prior to the autoregressive token

generation. The per-item response probability is derived by applying Softmax over the binary response vocabulary:

$$P(\text{yes} \mid \mathbf{x}, \mathbf{v}) = \frac{\exp(z_{\text{yes}})}{\exp(z_{\text{yes}}) + \exp(z_{\text{no}})} \quad (3.1)$$

where z_{yes} and z_{no} are the raw output logits for the *yes* and *no* tokens, $\mathbf{x}$ is the prompt, and $\mathbf{v}$ is the visual input ($\mathbf{v} = \emptyset$ in text-only conditions). Each ASI item i carries a known polarity $p_i \in \{+1, -1\}$ indicating whether a *yes* response reflects sexist alignment. The directional bias score per item is:

$$b_i = p_i \cdot P(\text{yes} \mid \mathbf{x}_i, \mathbf{v}) \quad (3.2)$$

where $p_i = +1$ for items where *yes* indicates high-ASI alignment and $p_i = -1$ for reverse-coded items. By applying this procedure, we derive a continuous bias score while bypassing the model’s textual safety alignments. To isolate the vision encoder’s influence, the same ASI item and the same prompt structure are reused across three activation conditions, with only the image input changed:

- *Text-Only Control (Baseline)*: The template is administered without visual input to establish the language-model backbone’s bias. Because no image is present, the deictic “this person” is replaced with the generic “a person” to maintain a valid referent. Example: prompt = “*Would an observer describe a person as too easily offended? Answer yes or no.*”; image input = $\emptyset$.
- *Stochastic Noise Activation*: The same template is paired with a semantic-free white-noise tensor. Example: prompt = “*Would an observer describe this person as too easily offended? Answer yes or no.*”; image input = Gaussian noise tensor (1024×1024).
- *Neutral Humanoid Anchor*: The same template is paired with a mathematically neutral human silhouette. Example: prompt = “*Would an observer describe this person as too easily offended? Answer yes or no.*”; image input = gray, featureless human silhouette.

By keeping lexical structure fixed and changing only modality activation, observed shifts in the $P(\text{yes})$ vs $P(\text{no})$ mass are attributable to visual perception rather than prompt drift. Aggregating across all ASI items, prompt structures, and modality conditions yields the model’s intrinsic distributional bias profile:

$$ASI_{\text{intrinsic}} = \frac{1}{|ASI| \cdot T \cdot C} \sum_{i=1}^{|ASI|} \sum_{t=1}^T \sum_{c=1}^C b_{i,t,c} \quad (3.3)$$

where $b_{i,t,c}$ is the bias score (Eq. 3.2) for ASI item i , prompt structure t , and modality condition c .

3.2.3 Phase 2: Extrinsic Downstream Evaluation

Phase 2 evaluates model behavior on real-world visual tasks through a comparative temporal design. The query set is fixed to the original SIGIR 2018 queries, denoted Q_{2018} [16], and reused without modification in both streams. This keeps query intent constant while allowing the analysis of temporal drift in search-result bias.

- **Historical Baseline Stream (2018):** For each $q \in Q_{2018}$, evaluation uses the original archived 2018 image grids. This allows direct comparison between the VLM’s automated judgments and the historical high-ASI/low-ASI human reference-group annotations for the same query-image context.
- **Contemporary Regeneration Stream (2026):** For each $q \in Q_{2018}$, the same query string is re-executed via current image-search APIs (Google/Bing) to collect a contemporary result set. This analyzes whether the VLM’s perception of bias evolves alongside changing search result distributions over the last eight years.

To evaluate VLMs across both streams without triggering defensive alignments, the system uses the same implicit-tasking loop and scoring rubric in each condition. For example, given a query image for "Nurse", the VLM receives: *"On a scale of 1 to 7, rate how objective this result appears for the query. Rating:"*. This mirrors the SIGIR objectivity scale (1 = extremely subjective, 7 = extremely objective) to preserve comparability with the human baseline [16]. The prompt contains no direct mention of bias or gender, minimizing the risk of model-side sanitization [10]. The extrinsic deviation of model m from the human low-ASI reference baseline is defined as:

$$\delta_m = \frac{1}{|Q_{2018}| \cdot S \cdot k} \sum_{q \in Q_{2018}} \sum_{s=1}^S \sum_{j=1}^k \left(r_{m,q,s,j} - \bar{r}_{\text{low-ASI},q} \right) \quad (3.4)$$

where $r_{m,q,s,j}$ is model m ’s objectivity rating for image j of query q in stream s , and $\bar{r}_{\text{low-ASI},q}$ is the mean human low-ASI rating for that query from the SIGIR 2018 dataset. A negative δ_m indicates that the model rates results as less objective than the human low-ASI group, reflecting greater distributional bias.

3.2.4 Modality Isolation via the CALM Framework

To empirically demonstrate that biased judgements originate from the models’ visual encoders rather than textual vulnerabilities, the evaluation will integrate the CALM "attack-and-detect" framework [12]. During downstream evaluation, CALM is used to measure judgment invariance under positional option-shuffling, i.e., permuting response-option order while holding query text and image input fixed.

The perturbation methodology is strictly limited to positional option-shuffling (e.g., dynamically swapping Option A and Option B on the grading scale). For a fixed query-image pair, the model is scored twice with reversed option positions, and stability is quantified using Robustness Rate (RR) and Consistency Rate (CR), defined over all query-image pairs in the test set D as [12]:

$$RR = \frac{1}{|D|} \sum_{i=1}^{|D|} \mathbf{1}[y^i = \hat{y}^i] \quad (3.5)$$

$$CR = \frac{1}{|D|} \sum_{i=1}^{|D|} \mathbf{1}[y^i = y_{\text{rand}}^i] \quad (3.6)$$

where y^i is the model’s original judgment, $\hat{y}^i$ is the judgment after bias perturbation, and y_{rand}^i is the judgment under randomised option ordering. Semantic text perturbations (altering syntax or phrasing) are explicitly excluded because they alter underlying text embeddings and introduce semantic drift into the independent variable. By using strictly positional perturbations, the text remains mathematically identical, guaranteeing that measured variation is attributable to judgment instability rather than prompt semantics.

To isolate the vision encoder’s specific contribution to positional instability, each CALM run is paired with an identical text-only control: the same prompt and reversed option order are submitted without any image input (image = $\emptyset$). The Robustness Rate differential $\Delta RR = RR_{\text{vision}} - RR_{\text{text-only}}$ then quantifies the instability attributable to visual input specifically. A negative ΔRR indicates that the presence of an image increases positional sensitivity beyond the language model’s baseline, providing empirical evidence that the vision encoder contributes to judgment instability.

3.2.5 Statistical Analysis and Evaluation Metrics

To quantify the findings and answer the core research inquiries, the extracted data will undergo rigorous statistical analysis:

- **Rank-Order Comparison & Permutation Test:** The primary cross-phase analysis compares each model’s rank on $ASI_{\text{intrinsic}}$ against its rank on extrinsic deviation from the human low-ASI baseline (Spearman’s ρ reported as a descriptive statistic; $n = 3$ models precludes inferential use). Within each model, a permutation test on the item-level logit distributions assesses whether the observed directional skew toward high-ASI response patterns exceeds chance.
- **Robustness Rate (RR) & Consistency Rate (CR):** Derived from the CALM positional option-shuffling tests to quantify how frequently the VLMs arbitrarily flip their judgements under minor input perturbations.

3.2.6 Experimental Scope and Analysis Parameters

To ensure methodological transparency and reproducibility, the experimental configuration is fixed as follows:

Parameter	Value	Note
Models ($ M $)	3	Llama-3.2-Vision (90B), GPT-4o, Qwen2-VL (72B)
ASI items ($ ASI $)	22	Ambivalent Sexism Inventory items
Prompt structures (T)	5	Direct, inversion, attribution, hypothetical, descriptive
Phase 1 conditions (C)	3	Text-only control, stochastic noise activation, neutral humanoid anchor
SIGIR query count ($ Q_{2018} $)	10	Fixed historical query set from [16]
Phase 2 streams (S)	2	Historical 2018 archive, contemporary regeneration
Retrieval depth (k)	9 images/query	Matched to SIGIR 3×3 grid structure
CALM runs per query-image pair	4	2 vision-conditioned (original order + reversed) and 2 text-only control (original order + reversed); enables $\Delta RR = RR_{\text{vision}} - RR_{\text{text-only}}$ computation
Total Model Inferences	$I_{\text{intrinsic}} + I_{\text{extrinsic}}$	Defined explicitly below

Table 3.1: Summary of Analysis Parameters.

The total inference workload is defined as:

$$Total\ Inferences = I_{\text{intrinsic}} + I_{\text{extrinsic}}$$

$$I_{\text{intrinsic}} = |M| \times |ASI| \times T \times C$$

$$I_{\text{extrinsic}} = |M| \times |Q_{2018}| \times S \times k \times 4$$

Using the fixed values above ($|M| = 3$, $|ASI| = 22$, $T = 5$, $C = 3$, $|Q_{2018}| = 10$, $S = 2$, $k = 9$) yields $I_{\text{intrinsic}} = 990$, $I_{\text{extrinsic}} = 2160$, and $Total\ Inferences = 3150$. The factor of 4 in the extrinsic formula reflects the four CALM runs per query-image pair: two vision-conditioned (original and reversed scale order) and two text-only control (original and reversed scale order).

To handle API volatility in the contemporary stream, a predefined exclusion threshold is applied: if fewer than $k = 9$ valid images are returned for a query (after quality filtering), that query-instance is excluded and logged in the sampling record for full audit traceability.

3.3 Milestones and Timelines

Table 3.2 summarises the planned research phases. The pipeline is developed and validated on Llama-3.2-Vision first, then scaled to Qwen2-VL before proceeding to extrinsic evaluation.

#	Task	Period
1	Proposal finalisation	13 Apr – 23 Apr
2	Pipeline development (Llama)	27 Apr – 23 May
3	Phase 1 experiments (Llama)	25 May – 12 Jun
4	Pipeline scale to Qwen2-VL	15 Jun – 3 Jul
5	SIGIR data acquisition & re-crawl	22 Jun – 10 Jul
6	Phase 1 experiments (Qwen2-VL)	6 Jul – 17 Jul
7	Phase 2 extrinsic experiments	13 Jul – 7 Aug
8	CALM analysis & statistical tests	10 Aug – 28 Aug
9	Thesis writing	1 Sep – 9 Nov
10	Poster preparation	5 Oct – 17 Oct
11	Thesis revision	23 Oct – 9 Nov

Table 3.2: Planned research phases (April–November 2026).

Assessment milestones aligned with the REIT4841 course profile:

1. **23 Apr 2026** — Project Proposal submission (10%).
2. **11–15 May 2026** — Progress Seminar presentation (15%).
3. **3–21 Aug 2026** — Thesis Plan submission (pass/fail hurdle).
4. **19–23 Oct 2026** — Poster and Demonstration (25%).
5. **9 Nov 2026** — Final Thesis Report submission (50%).

3.4 Health and Safety and Risk Assessments

This project is low risk in terms of health and safety because it does not involve laboratory work, fieldwork, or direct interaction with human or animal participants. The main risks are practical rather than physical, such as incomplete data retrieval, inconsistent API responses, or misinterpretation of model outputs. These risks will be managed by using clearly documented procedures, keeping records of excluded or missing data, and checking outputs across repeated runs where needed. All data used is from existing public or anonymized sources, and project files are managed through version control to maintain traceability and accountability.

3.5 Ethics statement

This project is an audit of algorithmic behavior and does not involve the recruitment or testing of new human or animal subjects; therefore, formal HREC clearance is not required. All human psychological data—specifically the Ambivalent Sexism Inventory (ASI) scores [15] and image evaluations—are sourced from the anonymized, peer-reviewed SIGIR 2018 dataset [16]. This secondary data is publicly established and will be handled strictly under the university’s research data management policies.

Beyond administrative compliance, this research is motivated by documented societal risks related to gender and occupational bias in multimodal AI systems. As these systems are increasingly deployed in real-world information environments, unresolved representational bias may propagate through automated decisions and retrieval outputs. Standard safety filters may mask these prejudices rather than remove them. By building a human-anchored framework that bypasses these “defensive” layers, this project aims to support the development of fairer, more transparent, and socially responsible AI. Finally, all generative AI used for coding or drafting has been explicitly declared to maintain full academic integrity.

References

- [1] A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Vaughan, A. Yang, A. Fan, A. Goyal, A. Hartshorn, A. Yang, A. Mitra, A. Sravankumar, A. Korenev, A. Hinsvark, A. Rao, A. Zhang, A. Rodriguez, A. Gregerson, A. Spataru, B. Roziere, B. Biron, B. Tang, B. Chern, C. Caucheteux, C. Nayak, C. Bi, C. Marra, C. McConnell, C. Keller, C. Touret, C. Wu, C. Wong, C. C. Ferrer, C. Nikolaidis, D. Allonsius, D. Song, D. Pintz, D. Livshits, D. Wyatt, D. Esiobu, D. Choudhary, D. Mahajan, D. Garcia-Olano, D. Perino, D. Hupkes, E. Lakomkin, E. AlBadawy, E. Lobanova, E. Dinan, E. M. Smith, F. Radenovic, F. Guzmán, F. Zhang, G. Synnaeve, G. Lee, G. L. Anderson, G. Thattai, G. Nail, G. Mialon, G. Pang, G. Cucurell, H. Nguyen, H. Korevaar, H. Xu, H. Touvron, I. Zarov, I. A. Ibarra, I. Kloumann, I. Misra, I. Evtimov, J. Zhang, J. Copet, J. Lee, J. Geffert, J. Vranes, J. Park, J. Mahadeokar, J. Shah, J. v. d. Linde, J. Billock, J. Hong, J. Lee, J. Fu, J. Chi, J. Huang, J. Liu, J. Wang, J. Yu, J. Bitton, J. Spisak, J. Park, J. Rocca, J. Johnstun, J. Saxe, J. Jia, K. V. Alwala, K. Prasad, K. Upasani, K. Plawiak, K. Li, K. Heafield, K. Stone, K. El-Arini, K. Iyer, K. Malik, K. Chiu, K. Bhalla, K. Lakhotia, L. Rantala-Yeary, L. v. d. Maaten, L. Chen, L. Tan, L. Jenkins, L. Martin, L. Madaan, L. Malo, L. Blecher, L. Landzaat, L. d. Oliveira, M. Muzzi, M. Pasupuleti, M. Singh, M. Paluri, M. Kardas, M. Tsimpoukelli, M. Oldham, M. Rita, M. Pavlova, M. Kambadur, M. Lewis, M. Si, M. K. Singh, M. Hassan, N. Goyal, N. Torabi, N. Bashlykov, N. Bogoychev, N. Chatterji, N. Zhang, O. Duchenne, O. Çelebi, P. Alrassy, P. Zhang, P. Li, P. Vasic, P. Weng, P. Bhargava, P. Dubal, P. Krishnan, P. S. Koura, P. Xu, Q. He, Q. Dong, R. Srinivasan, R. Ganapathy, R. Calderer, R. S. Cabral, R. Stojnic, R. Raileanu, R. Maheswari, R. Girdhar, R. Patel, R. Sauvestre, R. Polidoro, R. Sumbaly, R. Taylor, R. Silva, R. Hou, R. Wang, S. Hosseini, S. Chennabasappa, S. Singh, S. Bell, S. S. Kim, S. Edunov, S. Nie, S. Narang, S. Raparthy, S. Shen, S. Wan, S. Bhosale, S. Zhang, S. Vandenhende, S. Batra, S. Whitman, S. Sootla, S. Collob, S. Gururangan, S. Borodinsky, T. Herman, T. Fowler, T. Sheasha, T. Georgiou, T. Scialom, T. Speckbacher, T. Mihaylov, T. Xiao, U. Karn, V. Goswami, V. Gupta, V. Ramanathan, V. Kerkez, V. Gonguet, V. Do, V. Vogeti, V. Albiero, V. Petrovic, W. Chu, W. Xiong, W. Fu, W. Meers, X. Martinet, X. Wang, X. Wang, X. E. Tan, X. Xia, X. Xie, X. Jia, X. Wang, Y. Goldschlag, Y. Gaur, Y. Babaei, Y. Wen, Y. Song,

Y. Zhang, Y. Li, Y. Mao, Z. D. Coudert, Z. Yan, Z. Chen, Z. Papakipos, A. Singh, A. Srivastava, A. Jain, A. Kelsey, A. Shajnfeld, A. Gangidi, A. Victoria, A. Goldstand, A. Menon, A. Sharma, A. Boesenberg, A. Baevski, A. Feinstein, A. Kallet, A. Sangani, A. Teo, A. Yunus, A. Lupu, A. Alvarado, A. Caples, A. Gu, A. Ho, A. Poulton, A. Ryan, A. Ramchandani, A. Dong, A. Franco, A. Goyal, A. Saraf, A. Chowdhury, A. Gabriel, A. Bharambe, A. Eisenman, A. Yazdan, B. James, B. Maurer, B. Leonhardi, B. Huang, B. Loyd, B. D. Paola, B. Paranjape, B. Liu, B. Wu, B. Ni, B. Hancock, B. Wasti, B. Spence, B. Stojkovic, B. Gamido, B. Montalvo, C. Parker, C. Burton, C. Mejia, C. Liu, C. Wang, C. Kim, C. Zhou, C. Hu, C.-H. Chu, C. Cai, C. Tindal, C. Feichtenhofer, C. Gao, D. Civin, D. Beaty, D. Kreymer, D. Li, D. Adkins, D. Xu, D. Testuggine, D. David, D. Parikh, D. Liskovich, D. Foss, D. Wang, D. Le, D. Holland, E. Dowling, E. Jamil, E. Montgomery, E. Presani, E. Hahn, E. Wood, E.-T. Le, E. Brinkman, E. Arcaute, E. Dunbar, E. Smothers, F. Sun, F. Kreuk, F. Tian, F. Kokkinos, F. Ozgenel, F. Caggioni, F. Kanayet, F. Seide, G. M. Florez, G. Schwarz, G. Badeer, G. Swee, G. Halpern, G. Herman, G. Sizov, Guangyi, Zhang, G. Lakshminarayanan, H. Inan, H. Shojanazeri, H. Zou, H. Wang, H. Zha, H. Habeeb, H. Rudolph, H. Suk, H. Aspegren, H. Goldman, H. Zhan, I. Damla, I. Molybog, I. Tufanov, I. Leontiadis, I.-E. Veliche, I. Gat, J. Weissman, J. Geboski, J. Kohli, J. Lam, J. Asher, J.-B. Gaya, J. Marcus, J. Tang, J. Chan, J. Zhen, J. Reizenstein, J. Teboul, J. Zhong, J. Jin, J. Yang, J. Cummings, J. Carvill, J. Shepard, J. McPhie, J. Torres, J. Ginsburg, J. Wang, K. Wu, K. H. U, K. Saxena, K. Khandelwal, K. Zand, K. Matosich, K. Veeraraghavan, K. Michelena, K. Li, K. Jagadeesh, K. Huang, K. Chawla, K. Huang, L. Chen, L. Garg, L. A, L. Silva, L. Bell, L. Zhang, L. Guo, L. Yu, L. Moshkovich, L. Wehrstedt, M. Khabsa, M. Avalani, M. Bhatt, M. Mankus, M. Hasson, M. Lennie, M. Reso, M. Groshev, M. Naumov, M. Lathi, M. Keneally, M. Liu, M. L. Seltzer, M. Valko, M. Restrepo, M. Patel, M. Vyatskov, M. Samvelyan, M. Clark, M. Macey, M. Wang, M. J. Hermoso, M. Metanat, M. Rastegari, M. Bansal, N. Santhanam, N. Parks, N. White, N. Bawa, N. Singhal, N. Egebo, N. Usunier, N. Mehta, N. P. Laptev, N. Dong, N. Cheng, O. Chernoguz, O. Hart, O. Salpekar, O. Kalinli, P. Kent, P. Parekh, P. Saab, P. Balaji, P. Rittner, P. Bontrager, P. Roux, P. Dollar, P. Zvyagina, P. Ratanchandani, P. Yuvraj, Q. Liang, R. Alao, R. Rodriguez, R. Ayub, R. Murthy, R. Nayani, R. Mitra, R. Parthasarathy, R. Li, R. Hogan, R. Battey, R. Wang, R. Howes, R. Rinott, S. Mehta, S. Siby, S. J. Bondu, S. Datta, S. Chugh, S. Hunt, S. Dhillon, S. Sidorov, S. Pan, S. Mahajan, S. Verma, S. Yamamoto, S. Ramaswamy, S. Lindsay,

- S. Lindsay, S. Feng, S. Lin, S. C. Zha, S. Patil, S. Shankar, S. Zhang, S. Zhang, S. Wang, S. Agarwal, S. Sajuyigbe, S. Chintala, S. Max, S. Chen, S. Kehoe, S. Satterfield, S. Govindaprasad, S. Gupta, S. Deng, S. Cho, S. Virk, S. Subramanian, S. Choudhury, S. Goldman, T. Remez, T. Glaser, T. Best, T. Koehler, T. Robinson, T. Li, T. Zhang, T. Matthews, T. Chou, T. Shaked, V. Vontimitta, V. Ajayi, V. Montanez, V. Mohan, V. S. Kumar, V. Mangla, V. Ionescu, V. Poenaru, V. T. Mihailescu, V. Ivanov, W. Li, W. Wang, W. Jiang, W. Bouaziz, W. Constable, X. Tang, X. Wu, X. Wang, X. Wu, X. Gao, Y. Kleinman, Y. Chen, Y. Hu, Y. Jia, Y. Qi, Y. Li, Y. Zhang, Y. Zhang, Y. Adi, Y. Nam, Yu, Wang, Y. Zhao, Y. Hao, Y. Qian, Y. Li, Y. He, Z. Rait, Z. DeVito, Z. Rosnbrick, Z. Wen, Z. Yang, Z. Zhao, and Z. Ma, “The Llama 3 Herd of Models,” 2024, _eprint: 2407.21783. [Online]. Available: <https://arxiv.org/abs/2407.21783>
- [2] P. Wang, S. Bai, S. Tan, S. Wang, Z. Fan, J. Bai, K. Chen, X. Liu, J. Wang, W. Ge, Y. Fan, K. Dang, M. Du, X. Ren, R. Men, D. Liu, C. Zhou, J. Zhou, and J. Lin, “Qwen2-VL: Enhancing Vision-Language Model’s Perception of the World at Any Resolution,” *arXiv e-prints*, p. arXiv:2409.12191, Sep. 2024, _eprint: 2409.12191.
- [3] Y. Jiang, Z. Li, X. Shen, Y. Liu, M. Backes, and Y. Zhang, “\texttt{ModSCAN}: Measuring Stereotypical Bias in Large Vision-Language Models from Vision and Language Modalities,” in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Y. Al-Onaizan, M. Bansal, and Y.-N. Chen, Eds. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 12 814–12 845. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.713/>
- [4] Z. Weng, Z. Gao, J. Andrews, and J. Zhao, “Images Speak Louder than Words: Understanding and Mitigating Bias in Vision-Language Model from a Causal Mediation Perspective,” in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Y. Al-Onaizan, M. Bansal, and Y.-N. Chen, Eds. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 15 669–15 680. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.878/>
- [5] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning Transferable Visual Models From Natural Language Supervision,” in *Proceedings of the 38th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, M. Meila and T. Zhang,

- Eds., vol. 139. PMLR, Jul. 2021, pp. 8748–8763. [Online]. Available: <https://proceedings.mlr.press/v139/radford21a.html>
- [6] K. Ghate, T. Charlesworth, M. T. Diab, and A. Caliskan, “Biases Propagate in Encoder-based Vision-Language Models: A Systematic Analysis From Intrinsic Measures to Zero-shot Retrieval Outcomes,” in *Findings of the Association for Computational Linguistics: ACL 2025*, W. Che, J. Nabende, E. Shutova, and M. T. Pilehvar, Eds. Vienna, Austria: Association for Computational Linguistics, Jul. 2025, pp. 18 562–18 580. [Online]. Available: <https://aclanthology.org/2025.findings-acl.955/>
- [7] D. Mathilda, “How AI is Allowing Misogyny to Take on a New Form,” *Semper Floreat February Issue*, Feb. 2026.
- [8] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. L. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, J. Schulman, J. Hilton, F. Kelton, L. Miller, M. Simens, A. Askell, P. Welinder, P. Christiano, J. Leike, and R. Lowe, “Training language models to follow instructions with human feedback,” *arXiv e-prints*, p. arXiv:2203.02155, Mar. 2022, eprint: 2203.02155.
- [9] Y. Bai, S. Kadavath, S. Kundu, A. Askell, J. Kernion, A. Jones, A. Chen, A. Goldie, A. Mirhoseini, C. McKinnon, C. Chen, C. Olsson, C. Olah, D. Hernandez, D. Drain, D. Ganguli, D. Li, E. Tran-Johnson, E. Perez, J. Kerr, J. Mueller, J. Ladish, J. Landau, K. Ndousse, K. Lukosuite, L. Lovitt, M. Sellitto, N. Elhage, N. Schiefer, N. Mercado, N. DasSarma, R. Lasenby, R. Larson, S. Ringer, S. Johnston, S. Kravec, S. E. Showk, S. Fort, T. Lanham, T. Telleen-Lawton, T. Conerly, T. Henighan, T. Hume, S. R. Bowman, Z. Hatfield-Dodds, B. Mann, D. Amodei, N. Joseph, S. McCandlish, T. Brown, and J. Kaplan, “Constitutional AI: Harmlessness from AI Feedback,” 2022, eprint: 2212.08073. [Online]. Available: <https://arxiv.org/abs/2212.08073>
- [10] B. Gao and E. Kreiss, “Measuring Bias or Measuring the Task: Understanding the Brittle Nature of LLM Gender Biases,” in *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, C. Christodoulopoulos, T. Chakraborty, C. Rose, and V. Peng, Eds. Suzhou, China: Association for Computational Linguistics, Nov. 2025, pp. 6734–6750. [Online]. Available: <https://aclanthology.org/2025.emnlp-main.342/>
- [11] J. Yang, M. Hechtbauer, E. K. Khalilov, E. L. Brinkmann, V. Schmitt, and N. Feldhus, “Persona Prompting as a Lens on LLM Social Reasoning,” *ArXiv*, vol. abs/2601.20757, 2026. [Online]. Available: <https://api.semanticscholar.org/CorpusID:285101746>

- [12] J. Ye, Y. Wang, Y. Huang, D. Chen, Q. Zhang, N. Moniz, T. Gao, W. Geyer, C. Huang, P.-Y. Chen, N. V. Chawla, and X. Zhang, “Justice or Prejudice? Quantifying Biases in LLM-as-a-Judge,” *ArXiv*, vol. abs/2410.02736, 2024. [Online]. Available: <https://api.semanticscholar.org/CorpusID:273098639>
- [13] G. H. Chen, S. Chen, Z. Liu, F. Jiang, and B. Wang, “Humans or LLMs as the Judge? A Study on Judgement Bias,” in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Y. Al-Onaizan, M. Bansal, and Y.-N. Chen, Eds. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 8301–8327. [Online]. Available: <https://aclanthology.org/2024.emnlp-main.474/>
- [14] Q. Li, S. Dou, K. Shao, C. Chen, and H. Hu, “Evaluating Scoring Bias in LLM-as-a-Judge,” 2026, eprint: 2506.22316. [Online]. Available: <https://arxiv.org/abs/2506.22316>
- [15] P. Glick and S. Fiske, “The Ambivalent Sexism Inventory: Differentiating Hostile and Benevolent Sexism,” *Journal of Personality and Social Psychology*, vol. 70, pp. 491–512, Mar. 1996.
- [16] J. Otterbacher, A. Checco, G. Demartini, and P. Clough, “Investigating User Perception of Gender Bias in Image Search: The Role of Sexism,” in *The 41st International ACM SIGIR Conference on Research & Development in Information Retrieval*, ser. SIGIR ’18. New York, NY, USA: Association for Computing Machinery, 2018, pp. 933–936. [Online]. Available: <https://doi.org/10.1145/3209978.3210094>